

## 1. What is the vanishing gradient problem in deep neural networks? How does it affect training?

Ans- In the realm of deep learning, the optimization process plays a crucial role in training neural networks. Gradient descent, a fundamental optimization algorithm, can sometimes encounter two common issues: vanishing gradients and exploding gradients. In this article, we will delve into these challenges, providing insights into what they are, why they occur, and how to mitigate them. We will build and train a model, and learn how to face vanishing and exploding problems.

### What is Vanishing Gradient?

The vanishing gradient problem is a challenge that emerges during backpropagation when the derivatives or slopes of the activation functions become progressively smaller as we move backward through the layers of a neural network. This phenomenon is particularly prominent in deep networks with many layers, hindering the effective training of the model. The weight updates become extremely tiny, or even exponentially small, it can significantly prolong the training time, and in the worst-case scenario, it can halt the training process altogether.

### Why the Problem Occurs?

During backpropagation, the gradients propagate back through the layers of the network, they decrease significantly. This means that as they leave the output layer and return to the input layer, the gradients become progressively smaller. As a result, the weights associated with the initial levels, which accommodate these small gradients, are updated little or not at each iteration of the optimization process.

**The vanishing gradient problem** is particularly associated with the sigmoid and hyperbolic tangent (tanh) activation functions because their derivatives fall within the range of 0 to 0.25 and 0 to 1, respectively. Consequently, extreme weights become very small, causing the updated weights to closely resemble the original ones. This persistence of small updates contributes to the vanishing gradient issue.

The sigmoid and tanh functions limit the input values to the ranges  $[0,1]$  and  $[-1,1]$ , so that they saturate at 0 or 1 for sigmoid and -1 or 1 for Tanh. The derivatives at points becomes zero as they are moving. In these regions, especially when inputs are very small or large, the gradients are very close to zero. While this may not be a major concern in shallow networks with a few layers, it is a more pronounced issue in deep networks. When the inputs fall in saturated regions, the gradients approach zero, resulting in little update to the weights of the previous layer. In simple networks this does not pose much of a problem, but as more layers are added, these small gradients, which multiply between layers, decay significantly and consequently the first layer tears very slowly, and hinders overall model performance and can lead to convergence failure.

### How can we identify?

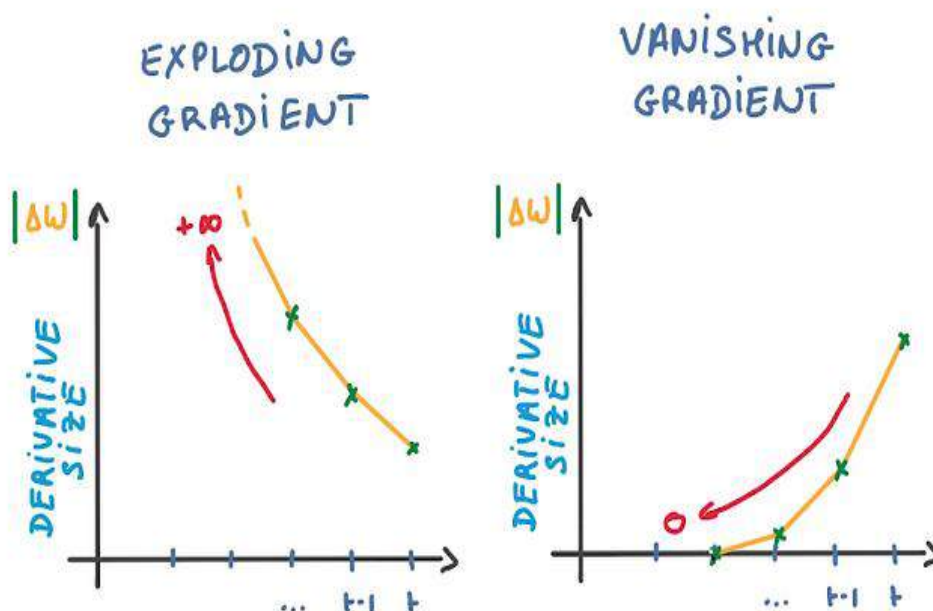
Identifying the vanishing gradient problem typically involves monitoring the training dynamics of a deep neural network.

- One key indicator is observing model weights **converging to 0** or stagnation in the improvement of the model's performance metrics over training epochs.
- During training, if the **loss function fails to decrease** significantly, or if there is erratic behaviour in the learning curves, it suggests that the gradients may be vanishing.

- Additionally, examining the gradients themselves during backpropagation can provide insights. **Visualization techniques**, such as gradient histograms or norms, can aid in assessing the distribution of gradients throughout the network.

**How can we solve the issue?**

- **Batch Normalization:** Batch normalization normalizes the inputs of each layer, reducing internal covariate shift. This can help stabilize and accelerate the training process, allowing for more consistent gradient flow.
- **Activation function:** Activation function like **Rectified Linear Unit (ReLU)** can be used. With **ReLU**, the gradient is 0 for negative and zero input, and it is 1 for positive input, which helps alleviate the vanishing gradient issue. Therefore, ReLU operates by replacing poor enter values with 0, and 1 for fine enter values, it preserves the input unchanged.
- **Skip Connections and Residual Networks (ResNets):** Skip connections, as seen in ResNets, allow the gradient to bypass certain layers during backpropagation. This facilitates the flow of information through the network, preventing gradients from vanishing.
- **Long Short-Term Memory Networks (LSTMs) and Gated Recurrent Units (GRUs):** In the context of recurrent neural networks (RNNs), architectures like LSTMs and GRUs are designed to address the vanishing gradient problem in sequences by incorporating gating mechanisms.
- **Gradient Clipping:** Gradient clipping involves imposing a threshold on the gradients during backpropagation. Limit the magnitude of gradients during backpropagation, this can prevent them from becoming too small or exploding, which can also hinder learning.



**What is Exploding Gradient?**

The exploding gradient problem is a challenge encountered during training deep neural networks. It occurs when the gradients of the network's loss function with respect to the weights (parameters) become excessively large.

### Why Exploding Gradient Occurs?

The issue of exploding gradients arises when, during backpropagation, the derivatives or slopes of the neural network's layers grow progressively larger as we move backward. This is essentially the opposite of the vanishing gradient problem.

The root cause of this problem lies in the weights of the network, rather than the choice of activation function. High weight values lead to correspondingly high derivatives, causing significant deviations in new weight values from the previous ones. As a result, the gradient fails to converge and can lead to the network oscillating around local minima, making it challenging to reach the global minimum point.

In summary, exploding gradients occur when weight values lead to excessively large derivatives, making convergence difficult and potentially preventing the neural network from effectively learning and optimizing its parameters.

As we discussed earlier, the update for the weights during backpropagation in a neural network is given by:

$$\Delta W_i = -\alpha \cdot \partial L / \partial W_i \quad \Delta W_i = -\alpha \cdot \partial W_i / \partial L$$

where,

- $\Delta W_i$  : The change in the weight  $W_i$
- $\alpha$ : The learning rate, a hyperparameter that controls the step size of the update.
- $L$ : The loss function that measures the error of the model.
- $\partial L / \partial W_i$ : The partial derivative of the loss function with respect to the weight  $W_i$ , which indicates the gradient of the loss function with respect to that weight.

The exploding gradient problem occurs when the gradients become very large during backpropagation. This is often the result of gradients greater than 1, leading to a rapid increase in values as you propagate them backward through the layers.

Mathematically, the update rule becomes problematic when  $|\nabla W_i| > 1$ , causing the weights to increase exponentially during training.

### How can we identify the problem?

Identifying the presence of exploding gradients in deep neural network requires careful observation and analysis during training. Here are some key indicators:

- The loss function exhibits erratic behaviour, oscillating wildly instead of steadily decreasing suggesting that the network weights are being updated excessively by large gradients, preventing smooth convergence.
- The training process encounters "NaN" (Not a Number) values in the loss function or other intermediate calculations.

- If network weights, during training exhibit significant and rapid increases in their values, it suggests the presence of exploding gradients.
- Tools like TensorBoard can be used to visualize the gradients flowing through the network.

#### How can we solve the issue?

- **Gradient Clipping:** It sets a maximum threshold for the magnitude of gradients during backpropagation. Any gradient exceeding the threshold is clipped to the threshold value, preventing it from growing unbounded.
- **Batch Normalization:** This technique normalizes the activations within each mini-batch, effectively scaling the gradients and reducing their variance. This helps prevent both vanishing and exploding gradients, improving stability and efficiency.

#### 2. Explain how Xavier initialization addresses the vanishing gradient problem.

Ans- Xavier initialization is a technique for initializing the weights of neural networks in a way that facilitates efficient training. It is named after Xavier Glorot, who introduced this method in a 2010 paper co-authored with Yoshua Bengio. In their influential research paper titled "Understanding the Challenges of Training Deep Feedforward Neural Networks," the authors conducted experiments to investigate a widely accepted rule of thumb in the field of deep learning. This rule involves initializing the weights of neural networks by selecting random values from a uniform distribution that ranges between -1 and 1. After this random initialization, the weights are then scaled down by a factor of 1 divided by the square root of the number of input units (denoted as 'n').

Xavier initialization aims to address the issue of maintaining variance in the forward and backward passes of a neural network, specifically when using certain activation functions like the hyperbolic tangent (tanh) and the logistic sigmoid. Regardless of how many input connections a neuron in a layer has, the variance of its output should be roughly the same. This property helps to prevent the vanishing or exploding gradient problem, which can occur if the variances change drastically between layers. Similarly, the variance of the gradients during backpropagation should also be roughly constant regardless of the number of neurons in the subsequent layer. This helps in maintaining stable training dynamics.

In Xavier initialization, the key factor is the number of inputs and outputs in the layers and not so much the method of randomization. The goal is to maintain the variance in bounds that enable effective learning with various activation functions.

##### Uniform Xavier Initialization

The Xavier initialization method is calculated as a random number with a uniform probability distribution (U) between the range  $-(1/\sqrt{n})$  and  $1/\sqrt{n}$ , where  $n$  is the number of inputs to the node:

$$\text{weight} = U \left[ -(1/\sqrt{n}), 1/\sqrt{n} \right]$$

##### Normal Xavier Initialization

This normalized Xavier initialization method is calculated as a random number with a uniform probability distribution (U) between the range  $-(\sqrt{6}/\sqrt{n + m})$  and  $\sqrt{6}/\sqrt{n + m}$ ,

where  $n$  is the number of inputs to the node (e.g. number of nodes in the previous layer) and  $m$  is the number of outputs from the layer (e.g. number of nodes in the current layer).

- $\text{weight} = U[-(\sqrt{6})/\sqrt{(n + m)}, \sqrt{6})/\sqrt{(n + m)}]$

By setting the standard deviation based on the number of inputs and outputs, it adjusts the scale of the weights in a way that keeps the network's activations within a reasonable range, regardless of the layer size.

The choice between Gaussian Xavier initialization and Uniform Xavier initialization may depend on the specific neural network architecture and the activation functions used.

## Applications

Xavier Initialization uses a factor of 2 in the numerator when initializing weights for activation functions like sigmoid and tanh because these functions have derivatives that are relatively small compared to the linear activation functions. Let's break down why this factor is used:

1. **Sigmoid Activation:** The sigmoid activation function squeezes its input into the range  $[0, 1]$ . For inputs that are far from zero, the derivative of the sigmoid function becomes very small, approaching zero. This means that during backpropagation, gradients can vanish when weights are initialized too large, making training difficult. The factor of 2 in Xavier Initialization helps counteract this effect by ensuring that the variance of the weights is appropriate to prevent the gradients from vanishing too quickly.
2. **Hyperbolic Tangent (tanh) Activation:** The tanh activation function also compresses its input into a range between -1 and 1. Similar to sigmoid, the tanh function has derivatives that become very small as inputs move away from zero. Therefore, without proper initialization, gradients can vanish during training. The factor of 2 in Xavier Initialization addresses this issue by controlling the variance of the weights, making it easier to train networks with tanh activation functions.

In both cases, the factor of 2 helps balance the initialization such that the weights are neither too small (leading to vanishing gradients) nor too large (leading to exploding gradients). Xavier Initialization aims to provide a suitable starting point for training by ensuring that the gradients are not too small for efficient learning.

### 3. What are some common activation functions that are prone to causing vanishing gradients?

Ans- The sigmoid function is one of the most popular activation functions used for developing deep neural networks. The use of sigmoid function restricted the training of deep neural networks because it caused the vanishing gradient problem. This caused the neural network to learn at a slower pace or in some cases no learning at all. This blog post aims to describe the vanishing gradient problem and explain how use of the sigmoid function resulted in it.


## Sigmoid function

Sigmoid functions are used frequently in neural networks to activate neurons. It is a logarithmic function with a characteristic S shape. The output value of the function is between 0 and 1. The


sigmoid function is used for activating the output layers in binary classification problems. It is calculated as follows:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

On the graph below you can see a comparison between the sigmoid function itself and its derivative. First derivatives of sigmoid functions are bell curves with values ranging from 0 to 0.25.




Our knowledge of how neural networks perform forward and backpropagation is essential to understanding the vanishing gradient problem.



### Forward Propagation

The basic structure of a neural network is an input layer, one or more hidden layers, and a single output layer. The weights of the network are randomly initialized during forward propagation. The input features are multiplied by the corresponding weights at each node of the hidden layer, and a bias is added to the net sum at each node. This value is then transformed into the output of the node using an activation function. To generate the output of the neural network, the hidden layer output is multiplied by the weights plus bias values, and the total is transformed using another activation function. This will be the predicted value of the neural network for a given input value.



## Back Propagation

As the network generates an output, the loss function( $C$ ) indicates how well it predicted the output. The network performs back propagation to minimize the loss. A back propagation method minimizes the loss function by adjusting the weights and biases of the neural network. In this method, the gradient of the loss function is calculated with respect to each weight in the network.

In back propagation, the new weight( $w_{\text{new}}$ ) of a node is calculated using the old weight( $w_{\text{old}}$ ) and product of the learning rate( $\eta$ ) and gradient of the loss function  $\left(\frac{\partial C}{\partial w}\right)$ .

$$W_{\text{new}} = W_{\text{old}} - \eta * \frac{\partial C}{\partial w}$$

With the chain rule of partial derivatives, we can represent gradient of the loss function as a product of gradients of all the activation functions of the nodes with respect to their weights. Therefore, the updated weights of nodes in the network depend on the gradients of the activation functions of each node.

For the nodes with sigmoid activation functions, we know that the partial derivative of the sigmoid function reaches a maximum value of 0.25. When there are more layers in the network, the value of the product of derivative decreases until at some point the partial derivative of the loss function approaches a value close to zero, and the partial derivative vanishes. We call this the vanishing gradient problem.


With shallow networks, sigmoid function can be used as the small value of gradient does not become an issue. When it comes to deep networks, the vanishing gradient could have a significant impact on performance. The weights of the network remain unchanged as the derivative vanishes. During back propagation, a neural network learns by updating its weights and biases to reduce the loss function. In a network with vanishing gradient, the weights cannot be updated, so the network cannot learn. The performance of the network will decrease as a result.

### Method to overcome the problem


The vanishing gradient problem is caused by the derivative of the activation function used to create the neural network. The simplest solution to the problem is to replace the activation function of the network. Instead of sigmoid, use an activation function such as ReLU.

Rectified Linear Units (ReLU) are activation functions that generate a positive linear output when they are applied to positive input values. If the input is negative, the function will return zero.





The derivative of a ReLU function is defined as 1 for inputs that are greater than zero and 0 for inputs that are negative. The graph shared below indicates the derivative of a ReLU function



If the ReLU function is used for activation in a neural network in place of a sigmoid function, the value of the partial derivative of the loss function will be having values of 0 or 1 which prevents the gradient from vanishing. The use of ReLU function thus prevents the gradient from vanishing. The problem with the use of ReLU is when the gradient has a value of 0. In such cases, the node is considered as a dead node since the old and new values of the weights remain the same. This situation can be avoided by the use of a leaky ReLU function which prevents the gradient from falling to the zero value.

Another technique to avoid the vanishing gradient problem is weight initialization. This is the process of assigning initial values to the weights in the neural network so that during back propagation, the weights never vanish.

In conclusion, the vanishing gradient problem arises from the nature of the partial derivative of the activation function used to create the neural network. The problem can be worse in deep neural networks using Sigmoid activation function. It can be significantly reduced by using activation functions like ReLU and leaky ReLU.

#### **4. Define the exploding gradient problem in deep neural networks. How does it impact training?**

Ans- The "exploding gradient problem" in deep neural networks occurs when the gradients used to update the network weights during backpropagation become excessively large, leading to unstable training where the model's parameters rapidly diverge from a stable solution, often causing the learning process to fail or oscillate wildly; essentially, the gradients "explode" and cause significant disruptions in weight updates, preventing the network from converging to an optimal solution.

Key points about exploding gradients:

- **Cause:**

This issue arises when gradients are multiplied across multiple layers in a deep network, and if these values are greater than 1, they can rapidly increase exponentially, causing the gradients to become very large.

- **Impact on training:**

- **Unstable learning:** Large gradient updates can cause weights to drastically change, making the loss function fluctuate wildly and preventing the model from converging to a good solution.
- **Numerical instability:** In extreme cases, exploding gradients can lead to numerical overflow, resulting in NaN values that cannot be used for further training.
- **Poor performance:** The model may struggle to learn effectively and exhibit erratic behaviour during training.

Factors contributing to exploding gradients:

- **Large initial weight values:**

When weights are initialized with large values, the gradient multiplication effect can be amplified, leading to exploding gradients.

- **Deep network architecture:**

Deeper networks have more layers where gradients can accumulate, increasing the likelihood of the exploding gradient problem.

- **Certain activation functions:**

Some activation functions like the vanilla sigmoid can contribute to exploding gradients when their gradients are not properly scaled.

How to mitigate exploding gradients:

- **Gradient clipping:**

A common technique where gradients are clipped to a certain maximum value, preventing them from becoming too large.

- **Weight initialization strategies:**

Using appropriate weight initialization methods like Xavier or He initialization can help stabilize gradients.

- **Adaptive learning rate optimizers:**

Optimizers like Adam or RMSprop can dynamically adjust learning rates to help control gradient updates.

- **Batch normalization:**

By normalizing activations within each layer, batch normalization can help mitigate exploding gradients.

## 5. What is the role of proper weight initialization in training deep neural networks?


Ans- While building and training neural networks, it is crucial to initialize the weights appropriately to ensure a model with high accuracy. If the weights are not correctly initialized, it may give rise to the Vanishing Gradient problem or the Exploding Gradient problem. Hence, selecting an appropriate weight initialization strategy is critical when training DL models. In this article, we will learn some of the most common weight initialization techniques, along with their implementation in Python using *Keras* in *TensorFlow*.

Following notations must be kept in mind while understanding the Weight Initialization Techniques. These notations may vary at different publications. However, the ones used here are the most common, usually found in research papers.

*fan\_in* = Number of input paths towards the neuron

*fan\_out* = Number of output paths towards the neuron

**Example:** Consider the following neuron as a part of a Deep Neural Network.



**For the above neuron,**

*fan\_in* = 3 (Number of input paths towards the neuron)

*fan\_out* = 2 (Number of output paths towards the neuron)

### Weight Initialization Techniques

#### Zero Initialization

As the name suggests, all the weights are assigned zero as the initial value is zero initialization. This kind of initialization is highly ineffective as neurons learn the same feature during each iteration. Rather, during any kind of constant initialization, the same issue happens to occur. Thus, constant initializations are not preferred.

In an attempt to overcome the shortcomings of Zero or Constant Initialization, **random initialization** assigns random values except for zeros as weights to neuron paths. However, assigning values

randomly to the weights, problems such as Overfitting, Vanishing Gradient Problem, Exploding Gradient Problem might occur.

**Random Initialization can be of two kinds:**

- Random Normal
- Random Uniform

**a) Random Normal:** The weights are initialized from values in a normal distribution.

**b) Random Uniform:** The weights are initialized from values in a uniform distribution.

$$w_i \sim N(0, 1)$$

In **Xavier/Glorot weight initialization**, the weights are assigned from values of a uniform distribution as follows:

$$w_i \sim U\left[-\sqrt{\frac{\sigma}{f_{an.in}+f_{an.out}}}, \sqrt{\frac{\sigma}{f_{an.in}+f_{an.out}}}\right]$$

Xavier/Glorot Initialization often termed as Xavier Uniform Initialization, is suitable for layers where the activation function used is **Sigmoid**.

In **Normalized Xavier/Glorot weight initialization**, the weights are assigned from values of a normal distribution as follows:

$$w_i \sim N(0, \sigma)$$

$$\sigma = \sqrt{\frac{6}{f_{an.in}+f_{an.out}}}$$

Xavier/Glorot Initialization, too, is suitable for layers where the activation function used is **Sigmoid**.

In **He Uniform weight initialization**, the weights are assigned from values of a uniform distribution as follows:

$$w_i \sim U\left[-\sqrt{\frac{6}{f_{an.in}}}, \sqrt{\frac{6}{f_{an.out}}}\right]$$

He Uniform Initialization is suitable for layers where **ReLU** activation function is used.

In **He Normal weight initialization**, the weights are assigned from values of a normal distribution as follows:

$$w_i \sim N[0, \sigma]$$

$$\sigma = \sqrt{\frac{2}{f_{an.in}}}$$

He Uniform Initialization, too, is suitable for layers where **ReLU** activation function is used.

Proper weight initialization in deep neural networks plays a crucial role by setting the starting values of the network's weights in a way that allows for efficient training, preventing issues like vanishing or exploding gradients, and ensuring the network can learn effectively by avoiding getting stuck in poor local minima during optimization; essentially, it determines the starting point for the learning process, significantly impacting convergence speed and overall model performance.

Key points about weight initialization:

- **Gradient stability:**

A well-initialized network helps maintain gradients within a reasonable range during backpropagation, preventing them from becoming too large (exploding) or too small (vanishing), which can hinder learning.

- **Activation function impact:**

Different activation functions require different weight initialization strategies to ensure appropriate activation values throughout the network.

- **Convergence speed:**

Proper initialization can lead to faster training by guiding the optimization process towards a good local minimum on the loss surface.

- **Common initialization methods:**

- **Xavier initialization:** Aims to maintain a balanced variance of activations across layers, particularly suitable for sigmoid and tanh activation functions.
- **He initialization:** Designed for ReLU activation, where weights are initialized with slightly larger values compared to Xavier.

What happens with poor weight initialization:

- **Slow learning:**

If weights are too small, gradients become very small, leading to slow learning progress.

- **Divergence:**

If weights are too large, gradients can explode, causing the network to become unstable and fail to converge.

## 6. Explain the concept of batch normalization and its impact on weight initialization techniques.

Ans- Batch normalization is a technique in deep learning that stabilizes the training process by normalizing the activations of a layer within each mini-batch, essentially making the input distribution to subsequent layers more consistent, which significantly reduces the sensitivity to initial weight values, allowing for less critical weight initialization strategies and potentially faster training with higher learning rates.

Key points about batch normalization:

- **How it works:**

For each mini-batch, the mean and standard deviation of the activations are calculated, and then each activation is normalized by subtracting the mean and dividing by the standard deviation.

- **Addressing "internal covariate shift":**

The primary benefit of batch normalization is mitigating the issue of "internal covariate shift," where the input distribution to later layers changes significantly during training due to weight updates, leading to instability.

Impact on weight initialization:

- **Less critical initialization:**

By stabilizing the input distributions, batch normalization reduces the need for carefully fine-tuned weight initialization techniques, as the network becomes less sensitive to the initial weight values.

- **Potential for larger initial weights:**

With batch normalization, you can often use larger initial weights without causing training instability, which can sometimes speed up the training process.


Important considerations:

- **Learnable parameters:**

Batch normalization introduces two additional learnable parameters per layer: a scaling factor (gamma) and a shift parameter (beta), allowing the network to adjust the normalized activations as needed.

- **Impact on generalization:**

While batch normalization often improves training stability, it can sometimes slightly decrease generalization ability due to the introduction of additional parameters.



## 7. Implement He initialization in Python using TensorFlow or PyTorch.

Ans-

[https://github.com/deep1185/PWSkills\\_assignments\\_1/blob/main/He\\_initialization\\_using\\_pytorch.ipynb](https://github.com/deep1185/PWSkills_assignments_1/blob/main/He_initialization_using_pytorch.ipynb)







